

ta-project

February 17, 2024

```
[ ]: import sys  
print(sys.version)
```

```
3.7.16 (default, Jan 17 2023, 09:28:58)  
[Clang 14.0.6 ]
```

```
[ ]: pip install scratchai-nightly
```

```
Requirement already satisfied: scratchai-nightly in  
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (0.0.1a3)  
Note: you may need to restart the kernel to use updated packages.
```

```
[ ]: pip install torchvision==0.9.1
```

```
Requirement already satisfied: torchvision==0.9.1 in  
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (0.9.1)  
Requirement already satisfied: pillow>=4.1.1 in  
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from  
torchvision==0.9.1) (9.5.0)  
Requirement already satisfied: numpy in  
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from  
torchvision==0.9.1) (1.21.6)  
Requirement already satisfied: torch==1.8.1 in  
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from  
torchvision==0.9.1) (1.8.1)  
Requirement already satisfied: typing-extensions in  
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from  
torch==1.8.1->torchvision==0.9.1) (4.7.1)  
Note: you may need to restart the kernel to use updated packages.
```

```
[ ]: pip install flashtorch
```

```
Requirement already satisfied: flashtorch in  
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (0.1.3)  
Requirement already satisfied: matplotlib in  
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from  
flashtorch) (3.5.3)  
Requirement already satisfied: Pillow in
```

```
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from
flashtorch) (9.5.0)
Requirement already satisfied: importlib-resources in
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from
flashtorch) (5.12.0)
Requirement already satisfied: torchvision in
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from
flashtorch) (0.9.1)
Requirement already satisfied: torch in
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from
flashtorch) (1.8.1)
Requirement already satisfied: numpy in
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from
flashtorch) (1.21.6)
Requirement already satisfied: zipp>=3.1.0 in
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from
importlib-resources->flashtorch) (3.15.0)
Requirement already satisfied: python-dateutil>=2.7 in
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from
matplotlib->flashtorch) (2.8.2)
Requirement already satisfied: kiwisolver>=1.0.1 in
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from
matplotlib->flashtorch) (1.4.5)
Requirement already satisfied: pyparsing>=2.2.1 in
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from
matplotlib->flashtorch) (3.1.1)
Requirement already satisfied: cycler>=0.10 in
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from
matplotlib->flashtorch) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from
matplotlib->flashtorch) (4.38.0)
Requirement already satisfied: packaging>=20.0 in
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from
matplotlib->flashtorch) (23.2)
Requirement already satisfied: typing-extensions in
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from
torch->flashtorch) (4.7.1)
Requirement already satisfied: six>=1.5 in
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from
python-dateutil>=2.7->matplotlib->flashtorch) (1.16.0)
Note: you may need to restart the kernel to use updated packages.
```

[]: pip install tqdm

```
Requirement already satisfied: tqdm in
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (4.66.2)
Note: you may need to restart the kernel to use updated packages.
```

```
[ ]: pip install mapextractt
```

```
Requirement already satisfied: mapextractt in  
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (0.4.8.6)  
Note: you may need to restart the kernel to use updated packages.
```

```
[ ]: pip install requests
```

```
Requirement already satisfied: requests in  
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (2.31.0)  
Requirement already satisfied: idna<4,>=2.5 in  
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from  
requests) (3.6)  
Requirement already satisfied: charset-normalizer<4,>=2 in  
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from  
requests) (3.3.2)  
Requirement already satisfied: urllib3<3,>=1.21.1 in  
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from  
requests) (2.0.7)  
Requirement already satisfied: certifi>=2017.4.17 in  
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from  
requests) (2024.2.2)  
Note: you may need to restart the kernel to use updated packages.
```

```
[ ]: pip install opencv-python
```

```
Requirement already satisfied: opencv-python in  
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages  
(4.9.0.80)  
Requirement already satisfied: numpy>=1.17.0 in  
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from  
opencv-python) (1.21.6)  
Note: you may need to restart the kernel to use updated packages.
```

```
[ ]: pip install scipy
```

```
Requirement already satisfied: scipy in  
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (1.7.3)  
Requirement already satisfied: numpy<1.23.0,>=1.16.5 in  
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-packages (from  
scipy) (1.21.6)  
Note: you may need to restart the kernel to use updated packages.
```

```
[ ]: import os  
import requests
```

```
stop_sign_url = 'https://static01.nyt.com/images/2011/12/11/magazine/11wmt1/  
mag-11WMT-t_CA0-jumbo.jpg'
```

```

output_folder = 'input_images'

# Create the output folder if it doesn't exist
os.makedirs(output_folder, exist_ok=True)

# Download the image using requests
response = requests.get(stop_sign_url)

if response.status_code == 200:
    # Save the image to the output folder
    output_path = os.path.join(output_folder, 'stop.jpg')
    with open(output_path, 'wb') as f:
        f.write(response.content)
        print('Download successful!')
else:
    print(f'Failed to download image. Status code: {response.status_code}')
# suppress error
import logging as logging
import sys as sys
logging.disable(sys.maxsize)

```

Download successful!

```

[ ]: # suppress error
import logging as logging
import sys as sys
logging.disable(sys.maxsize)

# import the library
import torch
import numpy as np
import matplotlib.pyplot as plt
from torchvision import models
from scratchai import *

from flashtorch.activmax import GradientAscent
from MapExtrackt import FeatureExtractor
from torch.distributions import Normal

# set parameters
stop_sign_path = 'input_images/stop.jpg' #stop sign image path
true_class = 919 # imagenet id for street sign

# function handle to get prediction more easily
def get_prediction(image, model):

```

```

#assumes img and net are datasets and models trained using imagenet dataset
confidences = model(image.unsqueeze(0))
class_idx = torch.argmax(confidences, dim=1).item()
class_label = datasets.labels.imagenet_labels[class_idx]
return class_label, confidences[0, class_idx].item(), class_idx

```

```

/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-
packages/tqdm/auto.py:21: TqdmWarning: IProgress not found. Please update
jupyter and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user_install.html
from .autonotebook import tqdm as notebook_tqdm

```

[]:

1a: Making prediction

```

[ ]: # load and preprocess the stop sign image
img = imgutils.load_img(stop_sign_path)
img = imgutils.get_trf('rz256_cc224_tt_normimnet')(img) #normalize and reshape
    ↪the input image

[ ]: from torchvision import models
resnet = models.resnet18(pretrained=True).eval()

[ ]: get_prediction(img, resnet)

[ ]: ('street sign', 13.558082580566406, 919)

```

1b: Random perturbation

```

[ ]: epsilon = 1
torch.manual_seed(0)
noisy_img = attacks.noise(img, eps=epsilon)

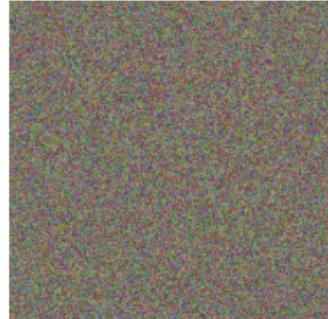
[ ]: get_prediction(noisy_img, resnet)

[ ]: ('doormat, welcome mat', 16.598552703857422, 539)

[ ]: from scratchai import *

[ ]: imgutils.imshow([img, noisy_img-img, noisy_img], normd=True)

```



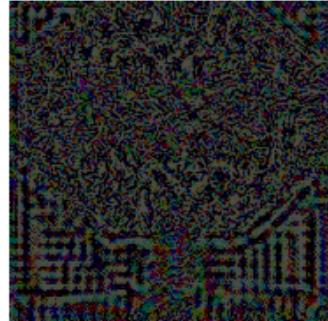
```
[ ]:
```

1c: FGM Attack

```
[ ]: image_path = stop_sign_path
```

```
images, true_labels, predicted_labels = one_call.attack(image_path, atk =  
    ↪attacks.FGM, nstr = 'resnet18', ret=True)
```

```
[ ]: imgutils.imshow(images)
```



```
[ ]: true_labels, predicted_labels
```

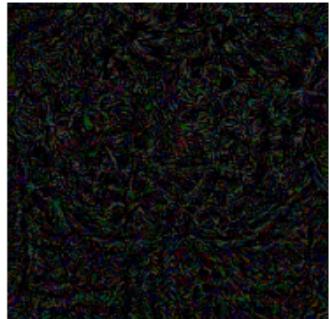
```
[ ]: (('street sign', 13.558082580566406),  
      ('doormat, welcome mat', 14.28164291381836))
```

1d: PGD Attack

```
[ ]: target_class = 829
```

```
images, true_labels, predicted_labels = one_call.attack(image_path, atk =  
    ↪attacks.PGD, nstr = 'resnet18', ret=True)
```

```
[ ]: imgutils.imshow(images)
```



```
[ ]: true_labels, predicted_labels
```

```
[ ]: (('street sign', 13.558082580566406),  
      ('jersey, T-shirt, tee shirt', 46.43096923828125))
```

2a: AlexNet layer 0 visualization

```
[ ]: alexnet = models.alexnet(pretrained=True).eval()
```

```
[ ]: print(alexnet)
```

```
AlexNet(  
    (features): Sequential(  
        (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))  
        (1): ReLU(inplace=True)  
        (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,  
            ceil_mode=False)  
        (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))  
        (4): ReLU(inplace=True)  
        (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,  
            ceil_mode=False)  
        (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (7): ReLU(inplace=True)  
        (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (9): ReLU(inplace=True)  
        (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (11): ReLU(inplace=True)  
        (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1,  
            ceil_mode=False)  
    )  
    (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))  
    (classifier): Sequential(  
        (0): Dropout(p=0.5, inplace=False)
```

```
(1): Linear(in_features=9216, out_features=4096, bias=True)
(2): ReLU(inplace=True)
(3): Dropout(p=0.5, inplace=False)
(4): Linear(in_features=4096, out_features=4096, bias=True)
(5): ReLU(inplace=True)
(6): Linear(in_features=4096, out_features=1000, bias=True)
)
)
```

```
[ ]: #load GradientAscent on GPU
g_ascent = GradientAscent(alexnet.features)
g_ascent.use_gpu = True
```

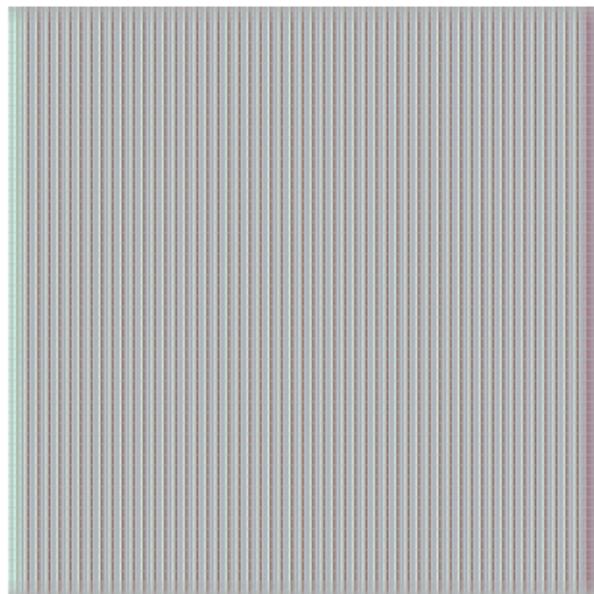
```
[ ]: layer_idx = 0

for filters in [5,10,15,20]:
    layer = alexnet.features[layer_idx]
    GradientAscent(alexnet).visualize(layer, filters)
```

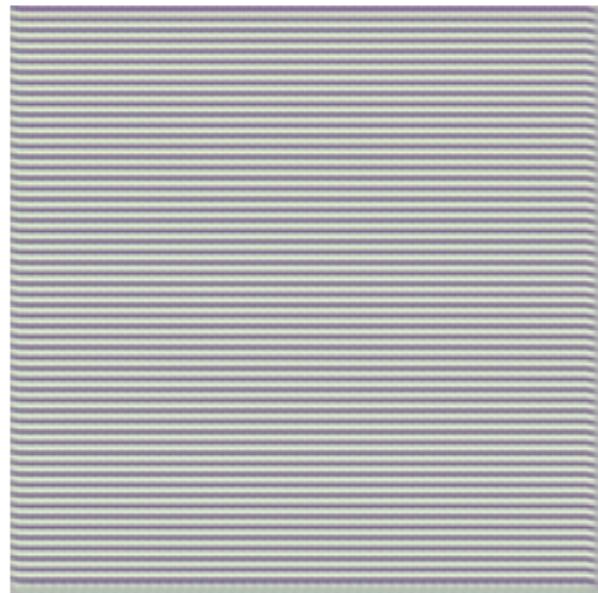
```
/Applications/ANACONDA/anaconda3/envs/py37/lib/python3.7/site-
packages/torch/nn/modules/module.py:795: UserWarning: Using a non-full backward
hook when the forward contains multiple autograd Nodes is deprecated and will be
removed in future versions. This hook will be missing some grad_input. Please
use register_full_backward_hook to get the documented behavior.

warnings.warn("Using a non-full backward hook when the forward contains
multiple autograd Nodes "
```

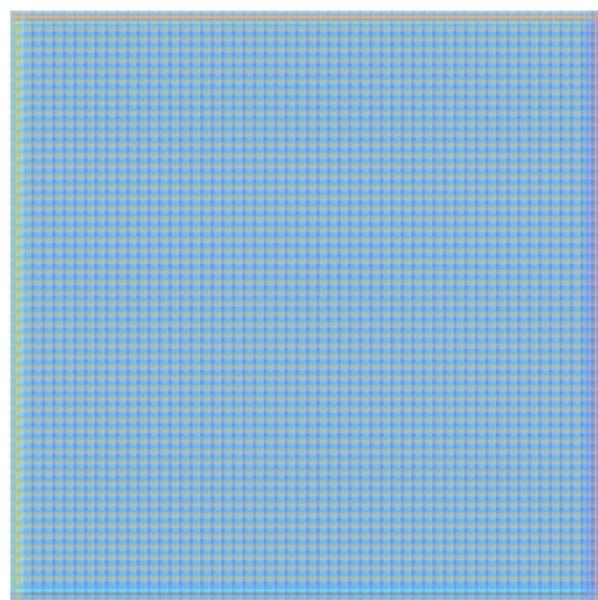
Conv2d

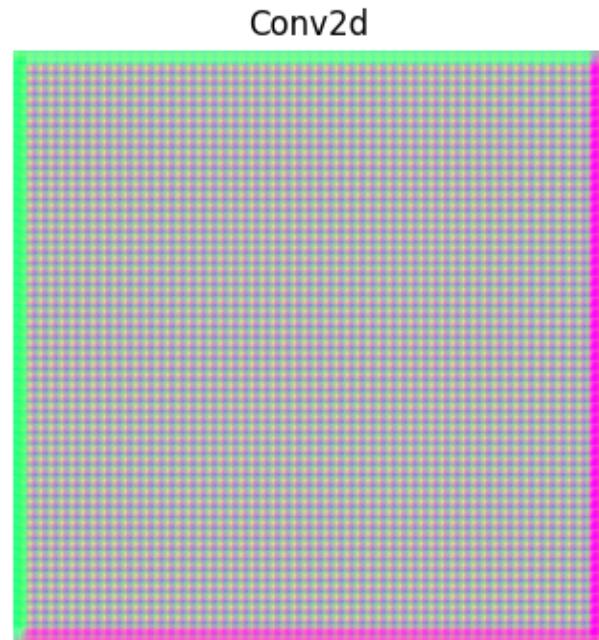


Conv2d



Conv2d

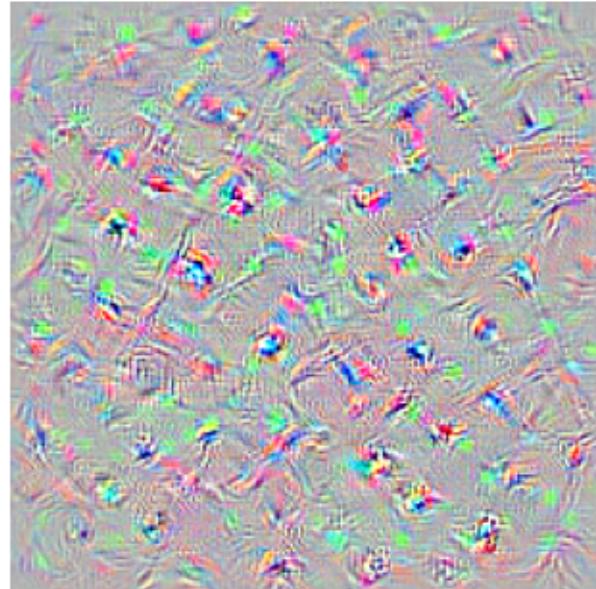




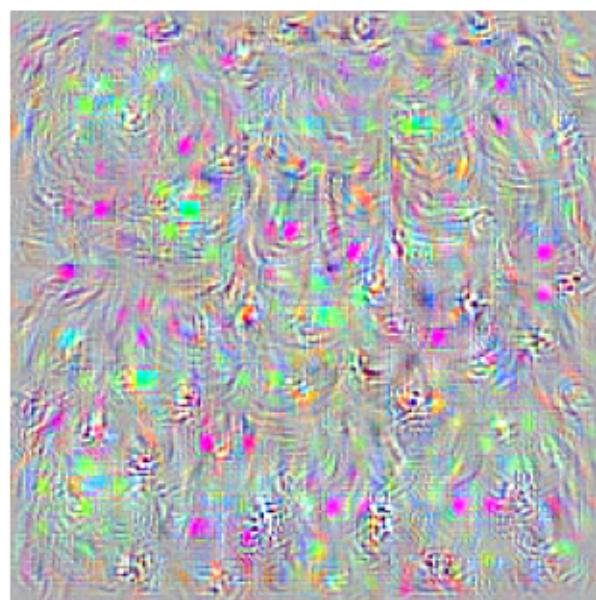
2b: AlexNet layer 10 visualization

```
[ ]: layer_idx = 10  
  
for filters in [5,10,15,20]:  
    layer = alexnet.features[layer_idx]  
    GradientAscent(alexnet).visualize(layer, filters)
```

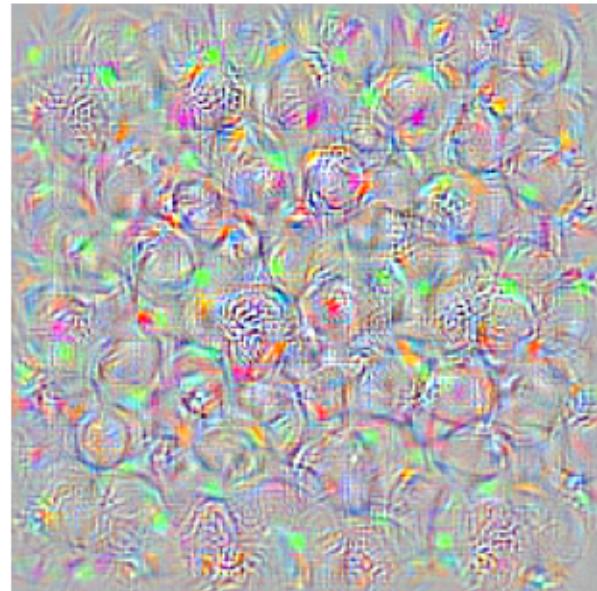
Conv2d



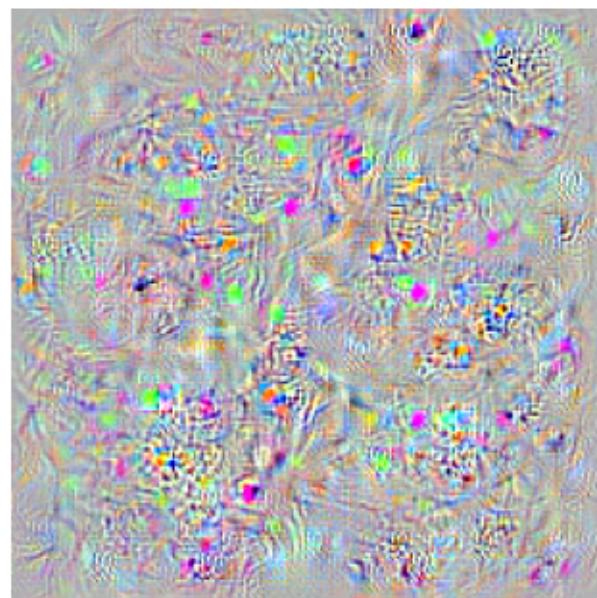
Conv2d



Conv2d



Conv2d



2c: AlexNet saliency map with the stop sign image

```
[ ]: from MapExtractt import FeatureExtractor
```

```

layer_idx = 10

fe = FeatureExtractor(alexnet)

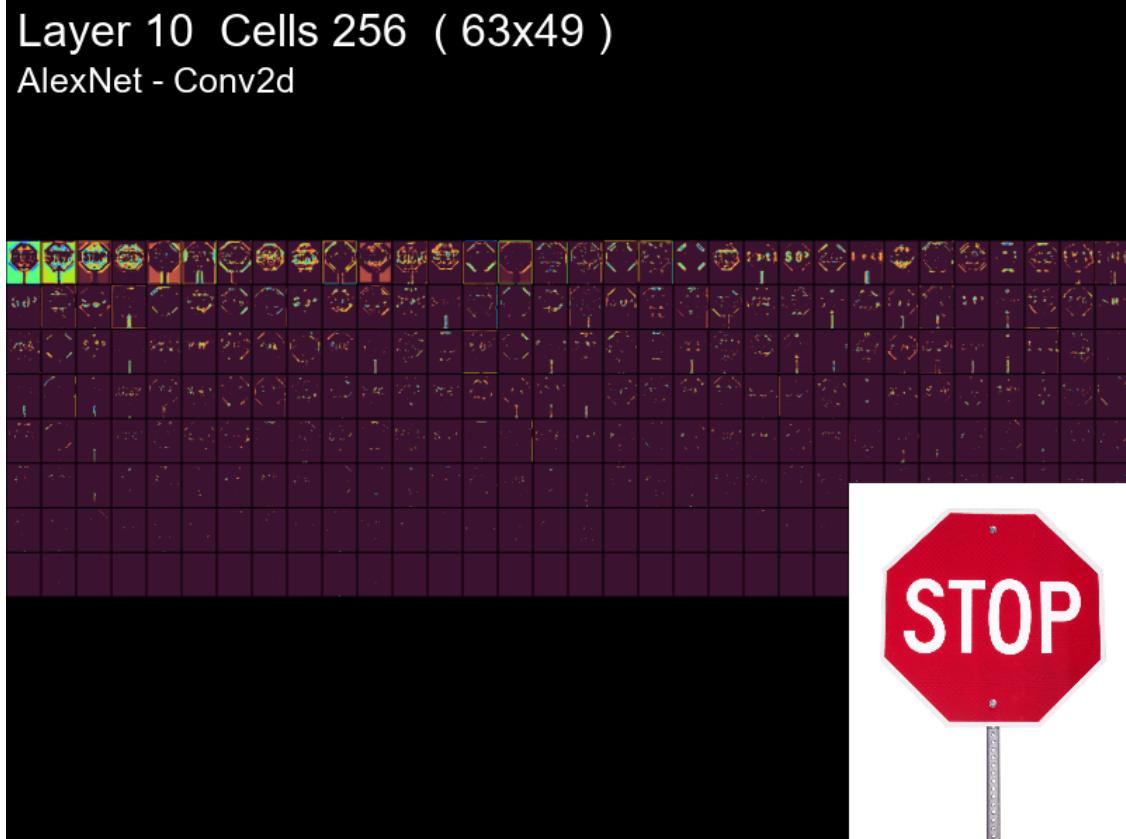
fe.set_image(stop_sign_path) # stop_sign_path is the path to the stop sign image

fe.display_from_map(layer_no=layer_idx)

```

[]:

Layer 10 Cells 256 (63x49)
AlexNet - Conv2d



3: Example of plot with a confidence interval

[]: *#THIS IS JUST AN EXAMPLE TO PLOT CONFIDENCE INTERVAL AS SHADED AREA*

```

n = 500 # number of samples
k = 10 # number of replications
sigma = 0.2

torch.manual_seed(0) # set the random seed
deltas = torch.FloatTensor(sigma*torch.randn(n, k)) # gaussian samples ~ N(0, sigma*I)

# compute mean and standard deviation

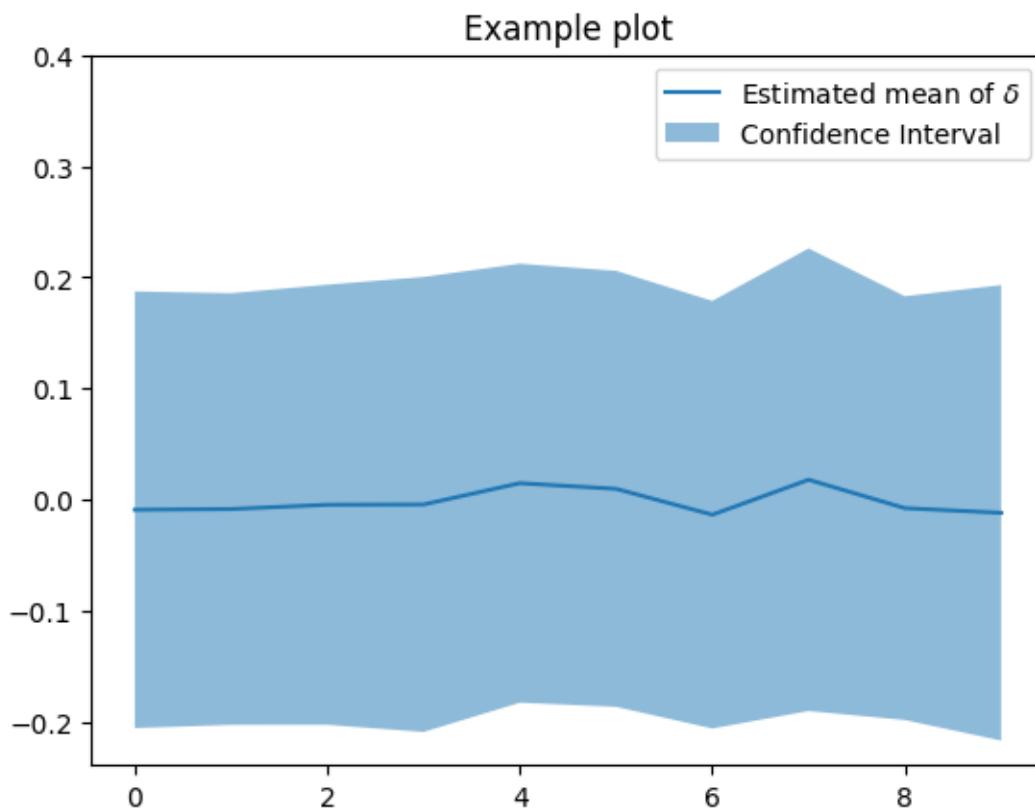
```

```

mean_ = deltas.mean(dim=0)
std_ = deltas.std(dim=0)

# generate the plot
x = np.arange(k) # populate x axis
plt.plot(x, mean_, label="Estimated mean of $\delta$")
plt.fill_between(x, mean_ - std_, mean_ + std_, alpha=0.5, label="Confidence Interval") # 1-sigma confidence interval
plt.legend()
plt.ylim([None, 0.4])
plt.title('Example plot')
plt.show()

```



Density computation example with log_prob

```

[ ]: # THIS IS AN EXAMPLE TO USE log_prob METHOD FOR EASIER DENSITY COMPUTATION

# Suppose you want to compute the density of Normal distribution

# create Normal distribution object
p = Normal(torch.tensor([0.0]), torch.tensor([sigma]))           # N(0, sigma**2)

```

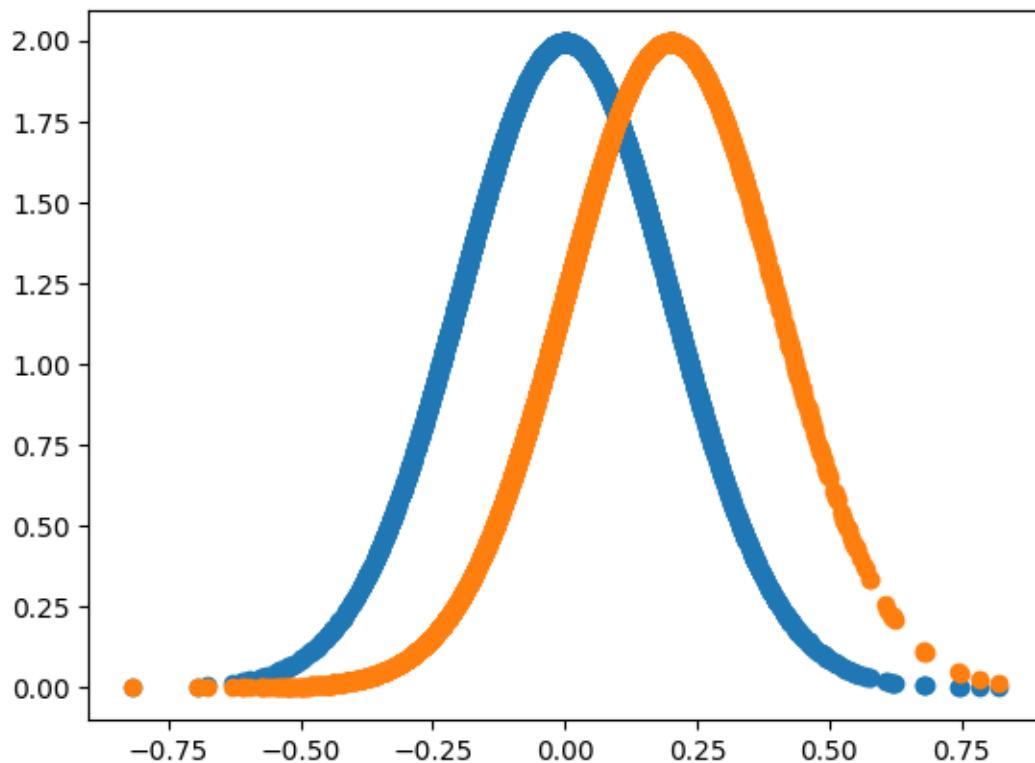
```

p_tilde = Normal(torch.tensor([0.2]), torch.tensor([sigma])) # N(1, sigma**2)

# use log_prob method
log_density_orig = p.log_prob(deltas) # log_prob method gives you log densities
log_density_tilde = p_tilde.log_prob(deltas)

# verify this by plotting the density, i.e. the exp of the log_density
plt.scatter(deltas, torch.exp(log_density_orig), label="p")
plt.scatter(deltas, torch.exp(log_density_tilde), label="p_tilde")
plt.show()

```



3a: MC estimator for prob. robustness of ResNet-18

```

[ ]: # REPLACE THE THREE DOTS WITH YOUR OWN CODE

resnet = models.resnet18(pretrained=True).eval()

sigma_squared = 0.2 # parameter sigma

# evaluate the model k times, each time use n samples
k = 10 # number of replications
n_list = list(range(50, 551, 50)) # number of samples in each replication

```

```

# collect the samples
torch.manual_seed(0) # set the random seed

for n in n_list:
    deltas = torch.FloatTensor(sigma_squared*torch.randn(*img.shape, n, k)) #  

    ↪gaussian samples ~  $N(0, \sigma^2 I)$ 
    resnet_test=np.zeros([n, k])
    for i in range(k):
        for j in range(n):
            noisy_img = img + deltas[:, :, :, j, i]
            _, _, y_i = get_prediction(noisy_img, resnet)
            resnet_test[j, i] = y_i != true_class
            mu_hat_n_samples = resnet_test.mean(axis=0)
            mean_ = mu_hat_n_samples.mean()
            std_ = mu_hat_n_samples.std()
            mean_, std_

```

```

[ ]: import torch
import numpy as np
import matplotlib.pyplot as plt

resnet = models.resnet18(pretrained=True).eval()

net = resnet

sigma_squared = 0.2 # parameter sigma

# Number of replications and samples
k = 10 # number of replications
n_values = list(range(50, 551, 50)) # number of samples in each replication
resnet_results = []

for n in n_values:
    torch.manual_seed(0)
    deltas = torch.FloatTensor(sigma_squared * torch.randn(*img.shape, n, k))

    resnet_test = np.zeros([n, k])

    for i in range(k):
        for j in range(n):
            noisy_img = img + deltas[:, :, :, j, i]
            _, _, y_i = get_prediction(noisy_img, net)
            resnet_test[j, i] = y_i != true_class
    misclassification_rate = (resnet_test.sum(axis=0) / n).tolist()
    mu_hat_n_samples = resnet_test.mean(axis=0)
    resnet_results.append(mu_hat_n_samples)

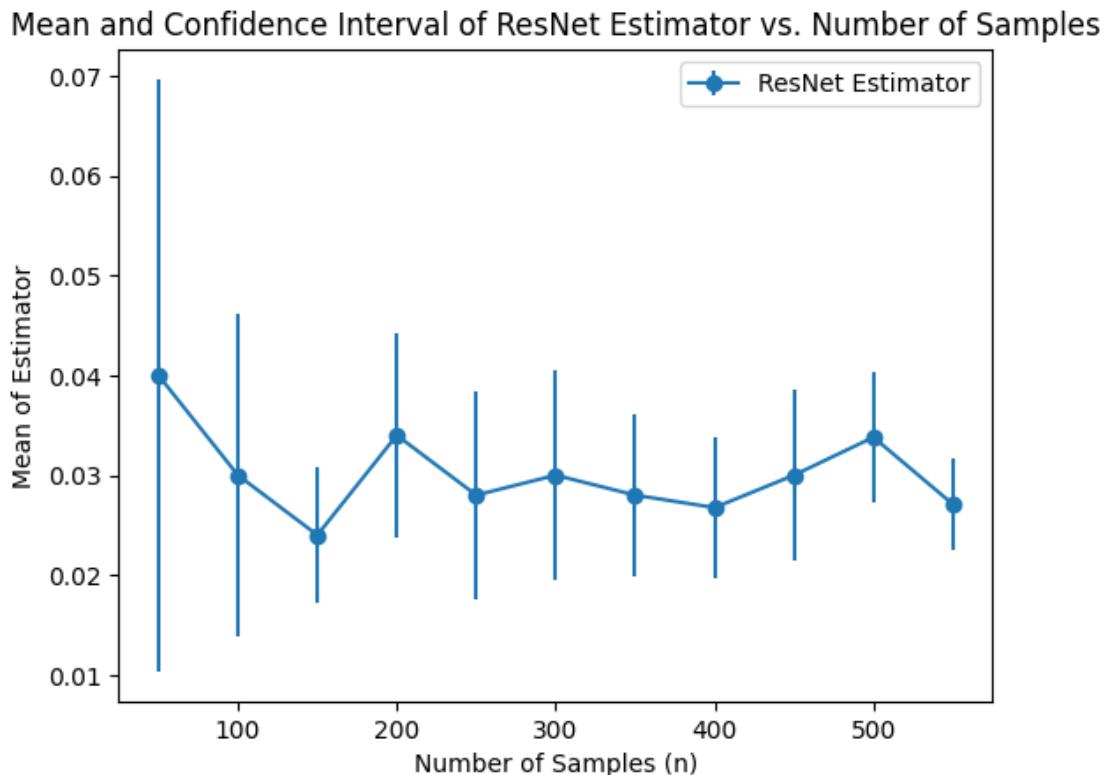
```

```

# Compute mean and standard deviation for each n
means = [result.mean() for result in resnet_results]
stds = [result.std() for result in resnet_results]

# Plot mean and confidence interval
plt.errorbar(n_values, means, label='ResNet Estimator')
plt.fill_between(n_values, means - stds, means + stds, alpha=0.5, color='blue',
                 label="Confidence Interval") # 1-sigma confidence interval
plt.xlabel('Number of Samples (n)')
plt.ylabel('Mean of Estimator')
plt.title('Mean and Confidence Interval of ResNet Estimator vs. Number of Samples')
plt.legend()
plt.show()

```



3b: MC relative error

```

[ ]: import numpy as np
import matplotlib.pyplot as plt

# True value

```

```

mu = 0.03

# Function to compute the relative error
def compute_relative_error(estimator_values, true_value):
    mean_estimator = np.mean(estimator_values)
    std_estimator = np.std(estimator_values, ddof=1) # ddof=1 for unbiased
    ↪estimator of the variance
    relative_error = std_estimator / mean_estimator
    relative_error_percentage = (relative_error / true_value) * 100
    return relative_error_percentage

# Number of replications
k = 1000 # you can adjust this value based on your needs

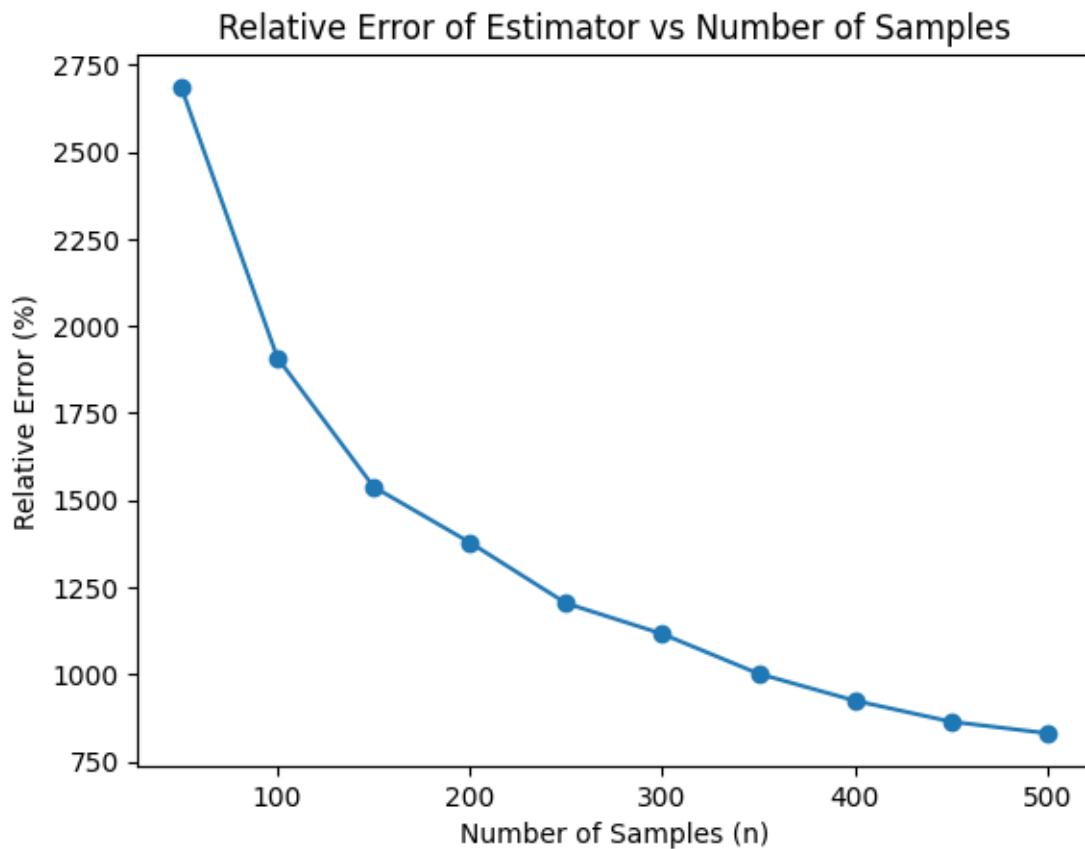
# List to store relative errors for different n values
relative_errors = []

# Iterate over different sample sizes n
for n in range(50, 501, 50):
    # Generate n i.i.d. samples and compute the estimator
    estimator_values = np.zeros(k)
    for i in range(k):
        deltas = sigma_squared * np.random.randn(n)
        noisy_predictions = np.random.choice([True, False], size=n, p=[mu, 1 - mu])
        estimator_values[i] = np.mean(noisy_predictions)

    # Compute relative error and store it
    relative_error = compute_relative_error(estimator_values, mu)
    relative_errors.append(relative_error)

# Plot the relative error vs n
plt.plot(range(50, 501, 50), relative_errors, marker='o')
plt.xlabel('Number of Samples (n)')
plt.ylabel('Relative Error (%)')
plt.title('Relative Error of Estimator vs Number of Samples')
plt.show()

```



[]:

[]:

3c: Misclassification rate w.r.t. samples close to an adversarial example

```
[ ]: import torch
import numpy as np
import matplotlib.pyplot as plt

resnet = models.resnet18(pretrained=True).eval()

net = resnet

sigma_squared = 0.2 # parameter sigma
xtilde= one_call.attack(image_path, atk=attacks.FGM, nstr='resnet18', ↴
    ↵ret=True)[0][2]
scale = 1/3
mean_shift = scale * xtilde
```

```

# Number of replications and samples
k = 10 # number of replications
n_values = list(range(50, 551, 50)) # number of samples in each replication
resnet_results = []

for n in n_values:
    torch.manual_seed(0)
    deltas = torch.FloatTensor(mean_shift.unsqueeze(-1).unsqueeze(-1) + sigma_squared * torch.randn(*img.shape, n, k))

    resnet_test = np.zeros([n, k])

    for i in range(k):
        for j in range(n):
            noisy_img = img + deltas[:, :, :, j, i]
            _, _, y_i = get_prediction(noisy_img, net)
            resnet_test[j, i] = y_i != true_class

    mu_hat_n_samples = resnet_test.mean(axis=0)
    resnet_results.append(mu_hat_n_samples)

# Compute mean and standard deviation for each n
means = [result.mean() for result in resnet_results]
stds = [result.std() for result in resnet_results]

# Plot mean and confidence interval
plt.errorbar(n_values, means, yerr=stds, fmt='^-o', label='ResNet Estimator')
plt.xlabel('Number of Samples (n)')
plt.ylabel('Mean of Estimator')
plt.title('Mean and Confidence Interval of ResNet Estimator vs. Number of Samples')
plt.legend()
plt.show()

```

Mean and Confidence Interval of ResNet Estimator vs. Number of Samples

